



太极图形课

第03讲 Advanced Data Layouts





太极图形课

第03讲 Advanced Data Layouts



Recap

- Metaprogramming
- Object-oriented programming



Reusability

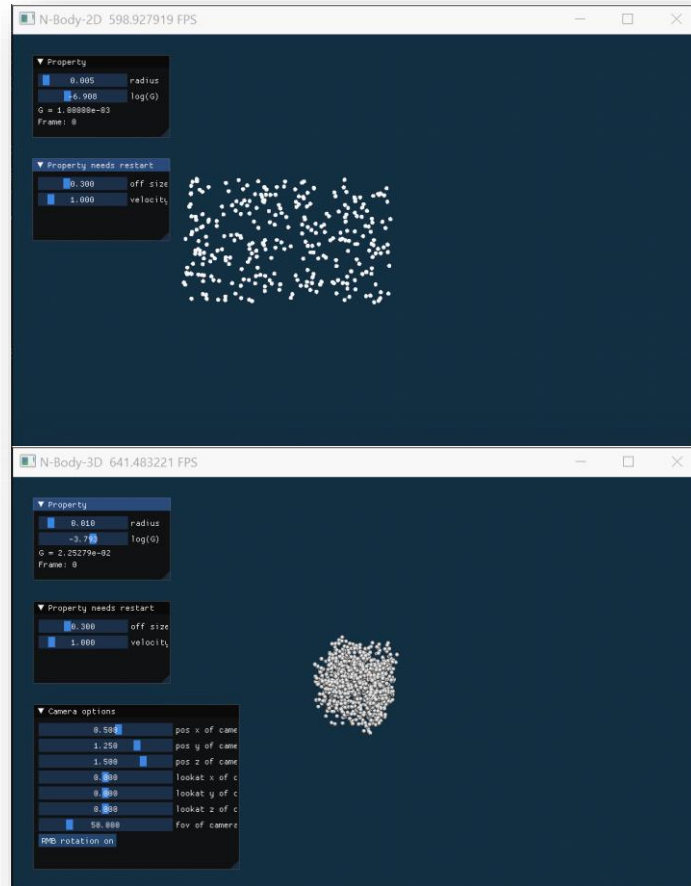


Extensibility




Maintainability

N-body systems




2/3D N-Body Dynamics
@Rabmelon



Solar System
@0xzhang


ODOP examples:




Diffraction @Y-jx007



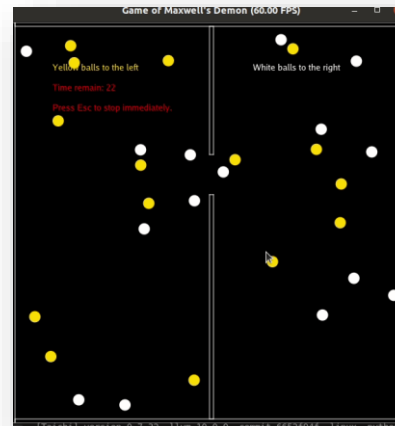
Bezier @Zydi



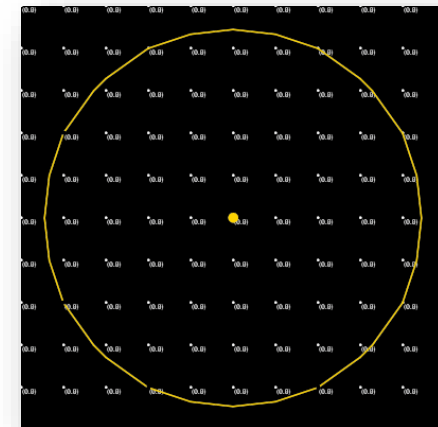
Moxi (墨戏) @Vineyo



Ant Colony
@theAfish




Maxwell's Demon game
@507C




Marching Squares
@AlbertLiDesign


Other HW assignments are welcome as well!



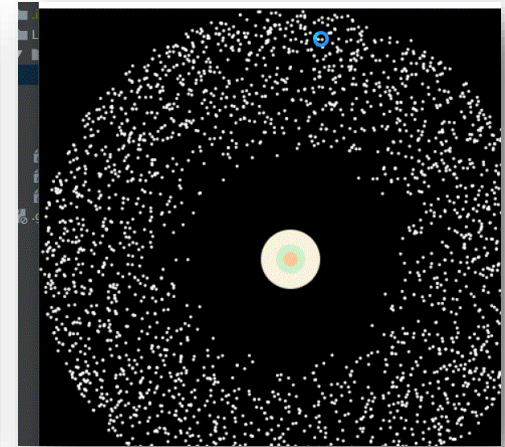
Mandelbulb
@rockeyshao



MPM88 + March Squares
@wangfeng70117



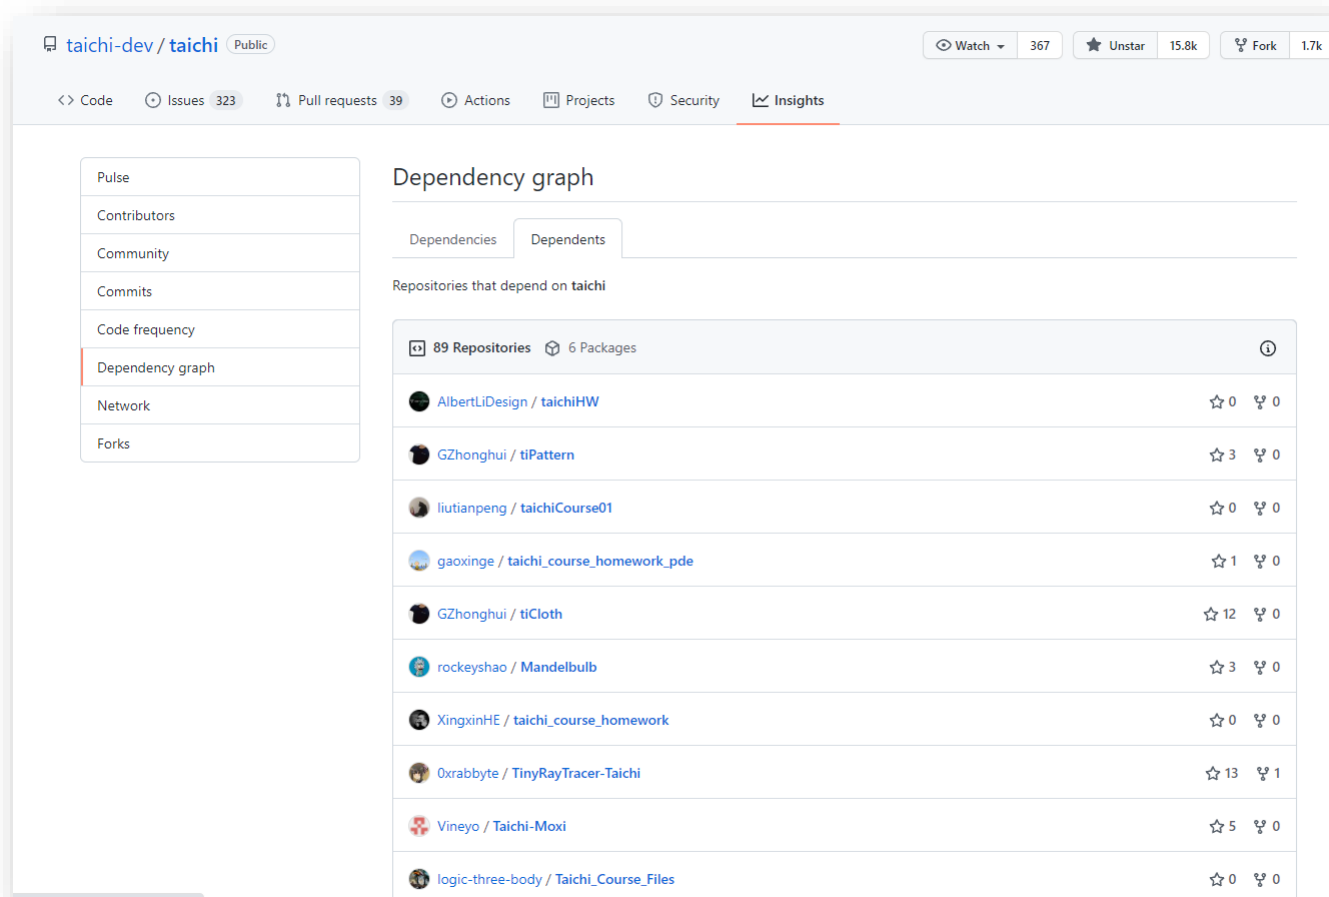
N-body with black hole(s)
@szl2



N-body with black hole(s)
@logic-three-body

Gifts for the gifted

- Check your Github issues 😊



The screenshot shows the Github repository page for `taichi-dev/taichi`. The repository is public and has 367 watchers, 15.8k stars, and 1.7k forks. The 'Issues' tab is selected, showing 323 issues. The 'Dependency graph' section is active, displaying a list of repositories that depend on `taichi`. The list includes 89 repositories and 6 packages. The following table summarizes the data from the screenshot:

Repository	Stars	Forks
AlbertLiDesign / taichiHW	0	0
GZhonghui / tiPattern	3	0
liutianpeng / taichiCourse01	0	0
gaoxing / taichi_course_homework_pde	1	0
GZhonghui / tiCloth	12	0
rockeyshao / Mandelbulb	3	0
XingxinHE / taichi_course_homework	0	0
0xrabyte / TinyRayTracer-Taichi	13	1
Vineyo / Taichi-Moxi	5	0
logic-three-body / Taichi_Course_Files	0	0



Outline Today

- Advanced dense data layouts
- Sparse data layouts

Outline Today

- Advanced dense data layouts
- Sparse data layouts



Performance

Performance

Performance


Advanced dense data layouts

Taichi

- `ti.field()`
- `@ti.kernel`
 - Optimized for `ti.field()`
- OOP
 - `@data_oriented`

Taichi: A **data-oriented** programming language

- ti.**field**()
- @ti.kernel
 - Optimized for ti.**field**()
- OOP
 - @**data_oriented**



Init

```
import taichi as ti

ti.init(ti.gpu)

# gravitational constant 6.67408e-11, using 1 for simplicity
G = 1
PI = 3.141592653

# number of planets
N = 300
# unit mass
m = 5
# galaxy size
galaxy_size = 0.4
# planet radius (for rendering)
planet_radius = 2
# init vel
init_vel = 120

# time-step size
h = 1e-5
# substepping
substepping = 10

# pos, vel and force of the planets
# Nx2 vectors
pos = ti.Vector.field(2, ti.f32, N)
vel = ti.Vector.field(2, ti.f32, N)
force = ti.Vector.field(2, ti.f32, N)
```

Data

```
@ti.kernel
def initialize():
    center = ti.Vector([0.5, 0.5])
    for i in range(N):
        theta = ti.random() * 4 * PI
        r = (ti.sqrt(ti.random())) * 0.7 + 0.3 * galaxy_size
        offset = r * ti.Vector([ti.cos(theta), ti.sin(theta)])
        pos[i] = center+offset
        vel[i] = [-offset.y, offset.x]
        vel[i] *= init_vel

@ti.kernel
def compute_force():
    # clear force
    for i in range(N):
        force[i] = ti.Vector([0.0, 0.0])

    # compute gravitational force
    for i in range(N):
        p = pos[i]
        for j in range(N):
            if i != j: # double the computation for a better memory footprint and load balance
                diff = p-pos[j]
                r = diff.norm(1e-5)

                # gravitational force -(G*m / r^2) * (diff/r) for i
                f = -G * m * m * (1.0/r)**3 * diff

                # assign to each particle
                force[i] += f

@ti.kernel
def update():
    dt = h/substepping
    for i in range(N):
        #symplectic euler
        vel[i] += dt*force[i]/m
        pos[i] += dt*vel[i]
```

Computation

```
gui = ti.GUI('N-body problem', (512, 512))

initialize()
while gui.running:



    for i in range(substepping):
        compute_force()
        update()

    ...

    gui.clear(0x12F41)
    gui.circles(pos.to_numpy(), color=0xffffffff, radius=planet_radius)
    gui.show()
```

Visualization


Performance @CPU...



Performance @GPU...

WALL CLOCK TIME

■ Computation ■ Data Access




better memory access!



搬砖 Example (a slide from @禹鹏)


```
def banzhuan():  
    for i in range(len(zhuan)):  
        move(zhuan[i])
```

```
@ti.kernel  
def banzhuan():  
    for i in zhuan:  
        move(zhuan[i])
```



.

..



.

..



Before we go: packed mode

- Initialized in *ti.init()*
- Decides whether to pad the data to the power of two
 - Default choice: *packed=False*, will do the padding
 - We assume *packed=True* in this class for simplicity

```
ti.init() # default: packed=False  
a = ti.field(ti.i32, shape=(18, 65)) # padded to (32, 128)
```

```
ti.init(packed=True)  
a = ti.field(ti.i32, shape=(18, 65)) # no padding
```

Taichi: optimized for data-access




```
x = ti.field(ti.i32, shape=16)

@ti.kernel
def fill():
    for i in x:
        x[i] = i

fill()
```

Taichi: optimized for data-access




```
x = ti.field(ti.i32, shape=16)
```

```
@ti.kernel
def fill():
    for i in x:
        x[i] = i
```

```
fill()
```

Taichi: optimized for data-access



Data in memory

Data prefetched

```
x = ti.field(ti.i32, shape=16)
```

```
@ti.kernel
```


```
def fill():
```

```
    for i in x:
```

```
        x[i] = i
```

```
fill()
```

Taichi: optimized for data-access





```
x = ti.field(ti.i32, shape=16)

@ti.kernel
def fill():
    for i in x:
        x[i] = i

fill()
```

Taichi: A **data-oriented** programming language





```
x = ti.field(ti.i32, shape=16)

@ti.kernel
def fill():
    for i in x:
        x[i] = i

fill()
```

Taichi: optimized for data-access




```
x = ti.field(ti.i32, shape=16)

@ti.kernel
def fill():
    for i in x:
        x[i] = i

fill()
```

Taichi: optimized for data-access


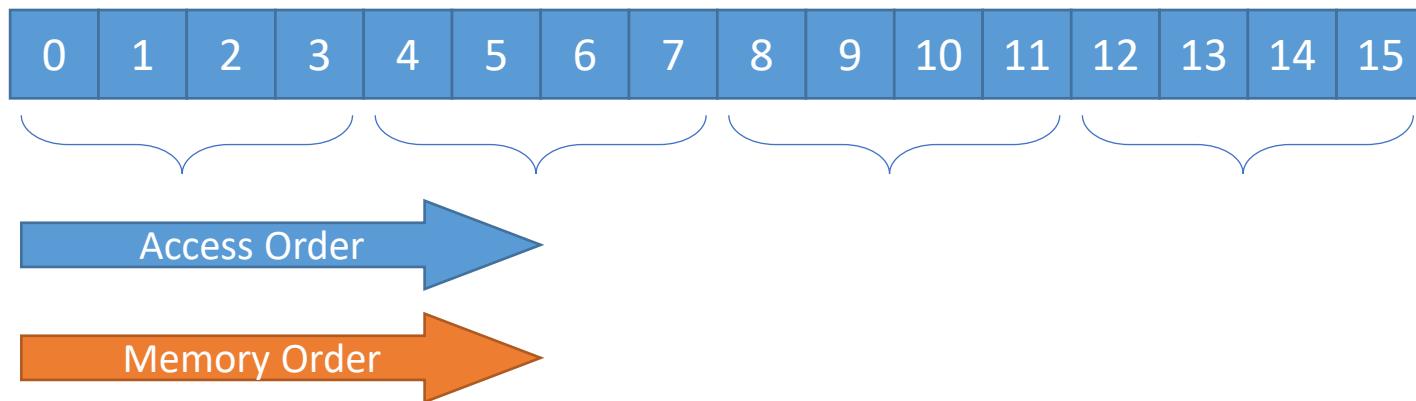


```
x = ti.field(ti.i32, shape=16)

@ti.kernel
def fill():
    for i in x:
        x[i] = i

fill()
```


Taichi: optimized for data-access




```
x = ti.field(ti.i32, shape=16)
```

```
@ti.kernel
def fill():
    for i in x:
        x[i] = i
```

```
fill()
```

How about multi-dimensional fields?

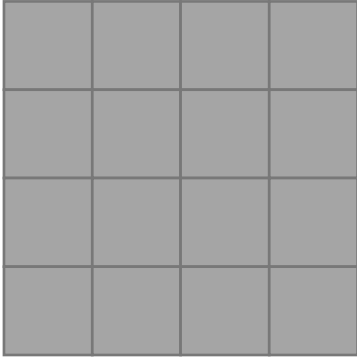


```
x = ti.field(ti.i32, shape = (4, 4))


@ti.kernel
def fill():
    for i,j in x:
        x[i,j] = 10*i + j

fill()
```

N-D fields are stored in our 1-D memory...



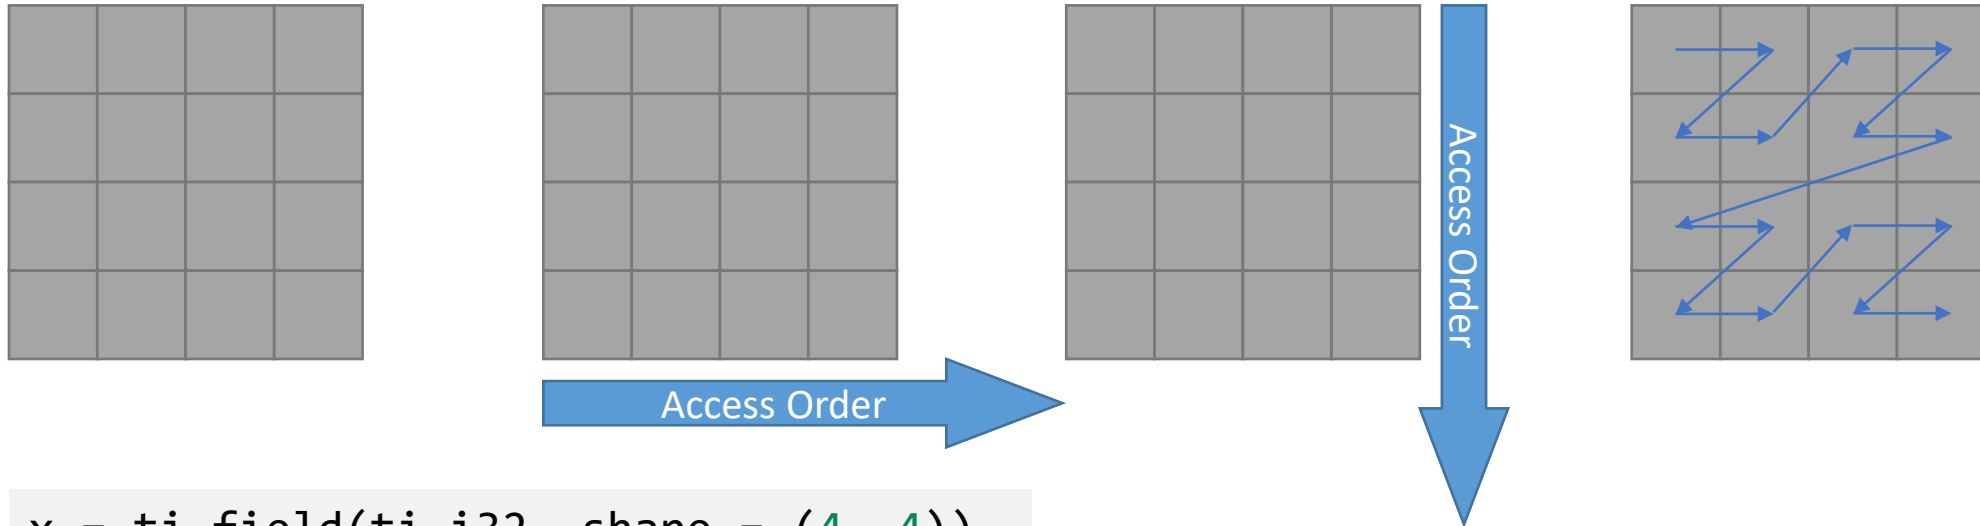
An N-D field
we think



An N-D field
we store



However the access pattern is not determined...




```
x = ti.field(ti.i32, shape = (4, 4))

@ti.kernel
def fill():
    for i,j in x:
        x[i,j] = 10*i + j

fill()
```

What we want:

- Store our data in a memory-access-friendly way.



Ideal memory layout of an N-D field:

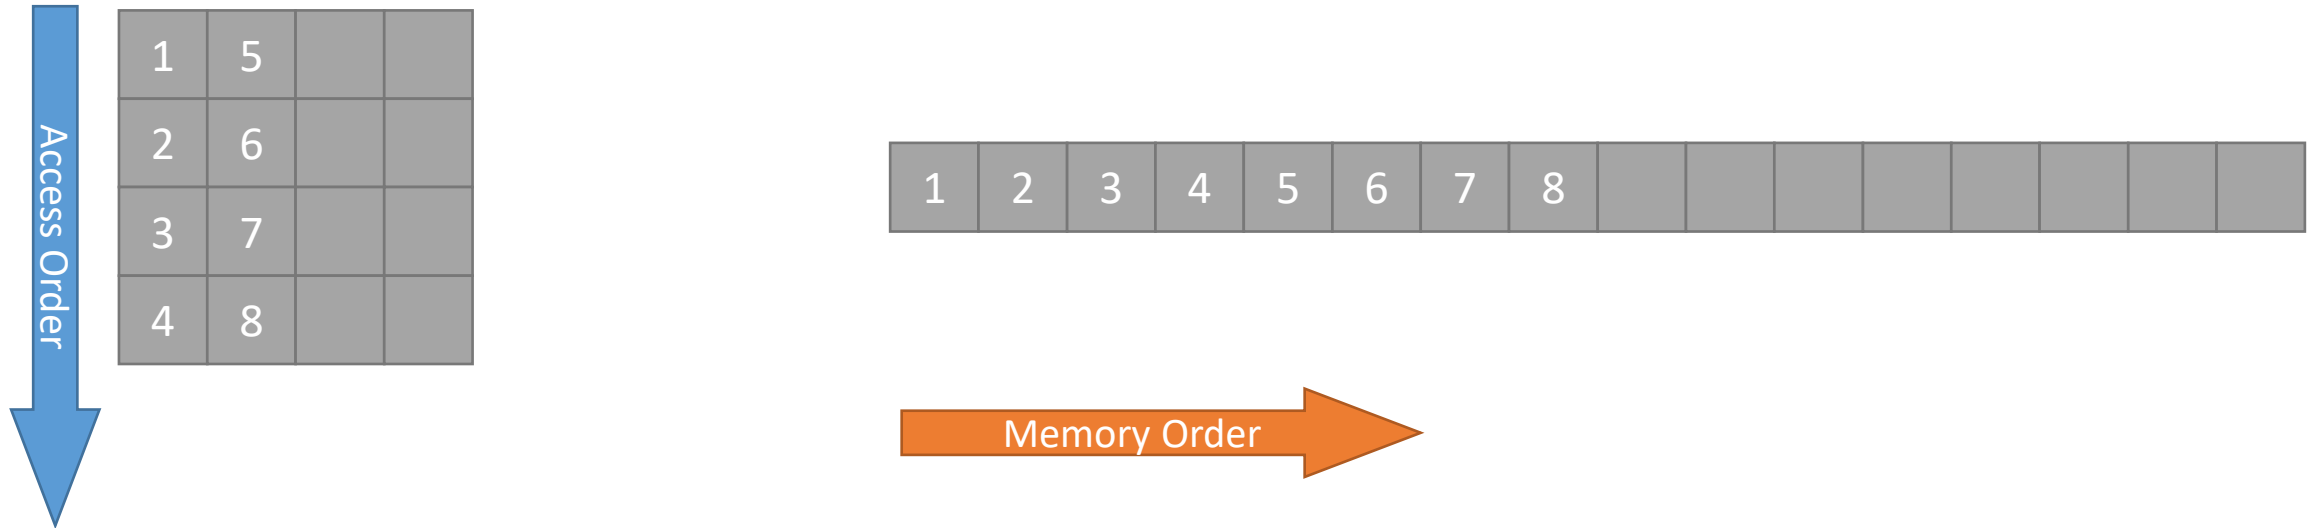
1	2	3	4
5	6	7	8

Access Order



1	2	3	4	5	6	7	8									
---	---	---	---	---	---	---	---	--	--	--	--	--	--	--	--	--

Memory Order

Ideal memory layout of an N-D field:





Ideal memory layout of an N-D field:



Memory Order

What we want:


- Store our data in a memory-access-friendly way.




What we want:

- memory-access-friendly?
- using `ti.field()`?



```
x = ti.field(ti.i32, shape = (4, 4))
```



Access Order



Access Order



Access row/col-major arrays/fields in C/C++


```
int x[3][2]; // row-major
int y[2][3]; // column-major

foo(){
    for (int i = 0; i < 3; i++) {
        for (int j = 0; j < 2; j++) {
            do_something(x[i][j]);
        }
    }

    for (int j = 0; j < 2; j++) {
        for (int i = 0; i < 3; i++) {
            do_something(y[j][i]);
        }
    }
}
```

C/C++

But that requires a huge stack in my brain ...



Upgrade your ti.field()



Layout 101: from *shape* to *ti.root*

```
x = ti.Vector.field(3, ti.f32, shape = 16)
```



```
x = ti.Vector.field(3, ti.f32)  
ti.root.dense(ti.i, 16).place(x)
```

ti.root In English:

Each cell of *root* has a *dense* container with *16 cells* along the *ti.i* axis. Each cell of a *dense* container has *field x*

Layout 101: from *shape* to *ti.root*

```
x = ti.Vector.field(3, ti.f32, shape = 16)
```



```
x = ti.Vector.field(3, ti.f32)  
ti.root.dense(ti.i, 16).place(x)
```

Root

ti.root In English:

Each cell of *root* has a *dense* container with 16 cells along the *ti.i* axis. Each cell of a *dense* container has *field* *x*

Layout 101: from *shape* to *ti.root*


```
x = ti.Vector.field(3, ti.f32, shape = 16)
```



```
x = ti.Vector.field(3, ti.f32)  
ti.root.dense(ti.i, 16).place(x)
```

ti.root In English:

Each cell of *root* has a *dense* container with *16 cells* along the *ti.i axis*. Each cell of a *dense* container has field *x*



Layout 101: from *shape* to *ti.root*


```
x = ti.Vector.field(3, ti.f32, shape = 16)
```



```
x = ti.Vector.field(3, ti.f32)  
ti.root.dense(ti.i, 16).place(x)
```

ti.root In English:

Each cell of *root* has a *dense* container with 16 cells along the *ti.i* axis. Each cell of a *dense* container has field *x*



ti.root: more examples:

```
x = ti.field(ti.f32, shape=())
```

```
x = ti.field(ti.f32, shape=3)
```

```
x = ti.field(ti.f32, shape=(3, 4))
```

```
x = ti.Matrix.field(2, 2, ti.f32, shape=5)
```

```
x = ti.field(ti.f32)  
ti.root.place(x)
```


```
x = ti.field(ti.f32)  
ti.root.dense(ti.i, 3).place(x)
```

```
x = ti.field(ti.f32)  
ti.root.dense(ti.ij, (3, 4)).place(x)
```

```
x = ti.Matrix.field(2, 2, ti.f32)  
ti.root.dense(ti.i, 5).place(x)
```


ti.root: the root of a *SNode-tree*

- SNode: Structural Node
- An SNode tree:
 - `ti.root` ← the root of the SNode-tree
 - `.dense()` ← a dense container describing shape
 - `.place(ti.field())` ← a field describing cell data
 - ...



ti.root: the root of a *SNode-tree*

- SNode: Structural Node
- An SNode tree:
 - `ti.root` ← the root of the SNode-tree
 - `.dense()` ← a dense container describing shape
 - `.place(ti.field())` ← a field describing cell data
 - ...



The SNode-tree


```
x = ti.field(ti.i32, shape = (4, 4))
```



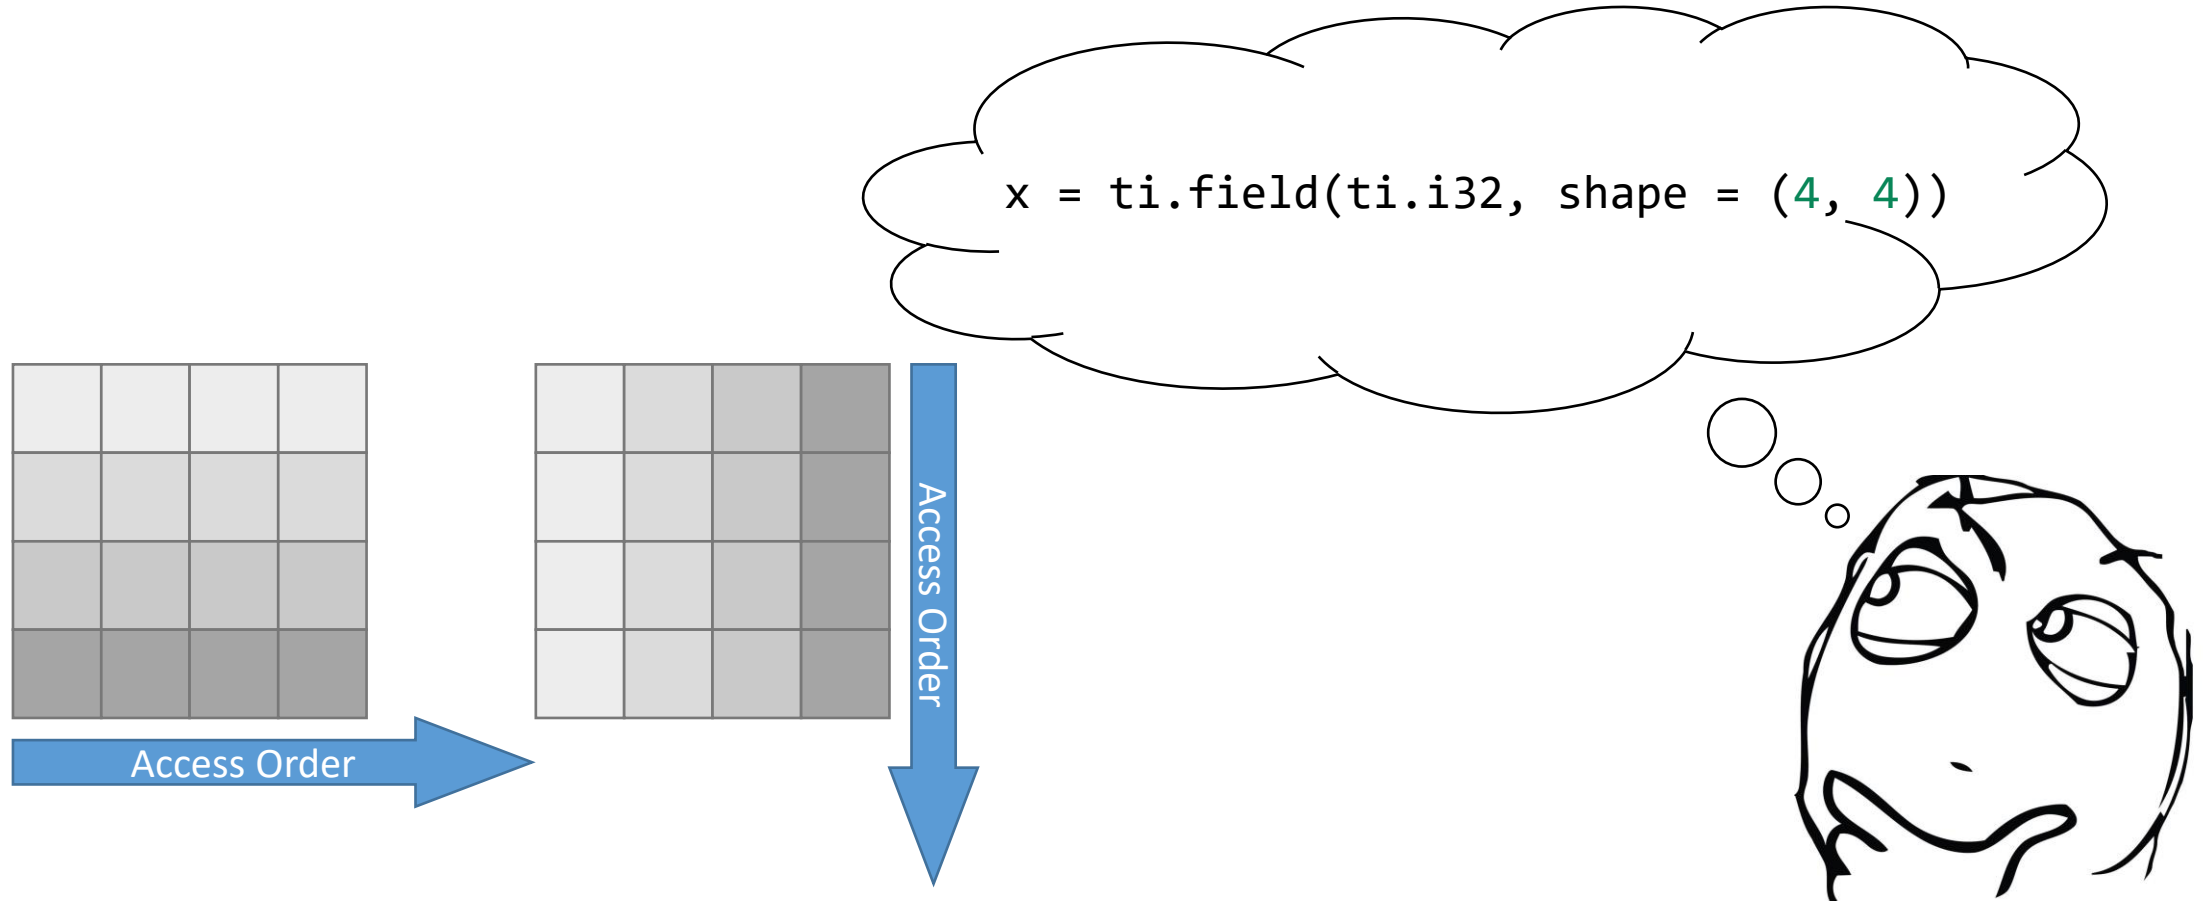
```
x = ti.field(ti.i32)  
ti.root.dense(ti.ij, (4, 4)).place(x)
```



```
x = ti.field(ti.i32)  
ti.root.dense(ti.i, 4).dense(ti.j, 4).place(x)
```




Row-major v.s. column-major




Row-major v.s. column-major

```
x = ti.field(ti.i32)
y = ti.field(ti.i32)
ti.root.dense(ti.i, 3).dense(ti.j, 2).place(x)    # row-major
ti.root.dense(ti.j, 2).dense(ti.i, 3).place(y)    # column-major
```





address: low High
x: x[0, 0] x[0, 1] x[1, 0] x[1, 1] x[2, 0] x[2, 1]
y: y[0, 0] y[1, 0] y[2, 0] y[0, 1] y[1, 1] y[2, 1]



Row-major access

```
x = ti.field(ti.i32)
ti.root.dense(ti.i, 4).dense(ti.j, 4).place(x)    # row-major
```



Access row/col-major fields

```
import taichi as ti
ti.init(arch = ti.cpu, cpu_max_num_threads=1)

x = ti.field(ti.i32)
ti.root.dense(ti.i, 3).dense(ti.j, 2).place(x)
# row-major

@ti.kernel
def fill():
    for i,j in x:
        x[i, j] = i*10 + j

@ti.kernel
def print_field():
    for i,j in x:
        print("x[" + i + ", " + j + "]=", x[i, j], sep='', end=' ')

fill()
print_field()
```

```
import taichi as ti
ti.init(arch = ti.cpu, cpu_max_num_threads=1)

x = ti.field(ti.i32)
ti.root.dense(ti.j, 2).dense(ti.i, 3).place(x)
# column-major

@ti.kernel
def fill():
    for i,j in x:
        x[i, j] = i*10 + j

@ti.kernel
def print_field():
    for i,j in x:
        print("x[" + i + ", " + j + "]=", x[i, j], sep='', end=' ')

fill()
print_field()
```


Access row/col-major fields

```
import taichi as ti
ti.init(arch = ti.cpu, cpu_max_num_threads=1)

x = ti.field(ti.i32)
ti.root.dense(ti.i, 3).dense(ti.j, 2).place(x)
# row-major

@ti.kernel
def fill():
    for i,j in x:
        x[i, j] = i*10 + j

@ti.kernel
def print_field():
    for i,j in x:
        print("x[" + i + ", " + j + "]=", x[i, j], sep='', end=' ')

fill()
print_field()
```

Loop over ***ti.j*** first

```
import taichi as ti
ti.init(arch = ti.cpu, cpu_max_num_threads=1)

x = ti.field(ti.i32)
ti.root.dense(ti.j, 2).dense(ti.i, 3).place(x)
# column-major

@ti.kernel
def fill():
    for i,j in x:
        x[i, j] = i*10 + j

@ti.kernel
def print_field():
    for i,j in x:
        print("x[" + i + ", " + j + "]=", x[i, j], sep='', end=' ')

fill()
print_field()
```

Loop over ***ti.i*** first

Access row/col-major arrays/fields in C/C++ v.s. in Taichi

```
int x[3][2]; // row-major
int y[2][3]; // column-major

foo(){
    for (int i = 0; i < 3; i++) {
        for (int j = 0; j < 2; j++) {
            do_something(x[i][j]);
        }
    }

    for (int j = 0; j < 2; j++) {
        for (int i = 0; i < 3; i++) {
            do_something(y[j][i]);
        }
    }
}
```

C/C++

```
x = ti.field(ti.i32)
y = ti.field(ti.i32)
ti.root.dense(ti.i, 3).dense(ti.j, 2).place(x) # row-major
ti.root.dense(ti.j, 2).dense(ti.i, 3).place(y) # column-major


@ti.kernel
def foo():
    for i,j in x:
        do_something(x[i, j])

    for i,j in y:
        do_something(y[i, j])
```

Taichi (Python)


A special case:

```
x = ti.field(ti.i32)  
ti.root.dense(ti.i, 4).dense(ti.i, 4).place(x)    # what is this?
```



Hierarchical layouts

```
x = ti.field(ti.i32)  
ti.root.dense(ti.i, 4).dense(ti.i, 4).place(x)    # A hierarchical 1-D field
```



Access a hierarchical 1-D field

```
import taichi as ti
ti.init(arch = ti.cpu)

x = ti.field(ti.i32)
ti.root.dense(ti.i, 4).dense(ti.i, 4).place(x)
# A hierarchical 1-D field

@ti.kernel
def print_id():
    for i in x:
        print(i, end = ' ')

print_id()
```

My first execution:

[Taichi] Starting on arch=x64

0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15

My second execution:

[Taichi] Starting on arch=x64

0 1 2 3 4 5 6 7 12 13 14 15 8 9 10 11


My third execution:

[Taichi] Starting on arch=x64

4 5 6 7 12 13 14 15 0 1 2 3 8 9 10 11


Hierarchical layouts

```
x = ti.field(ti.i32)
ti.root.dense(ti.i, 4).dense(ti.i, 4).place(x)    # A hierarchical 1-D field
```




- Access like a 1-D field
- Store like a 2-D field (in blocks)

Block-major access?





```
x = ti.field(ti.i32, shape = (4, 4))
```



Block-major access using hierarchical fields

```
x = ti.field(ti.i32)  
ti.root.dense(ti.ij, (2,2)).dense(ti.ij, (2,2)).place(x)    # block-major
```



Flat layouts v.s. hierarchical layouts

```
import taichi as ti
ti.init(arch = ti.cpu, cpu_max_num_threads=1)

z = ti.field(ti.i32, shape=(4,4))
# a row-major flat layout, size = 4x4

@ti.kernel
def fill():
    for i,j in z:
        z[i, j] = i*10 + j

@ti.kernel
def print_field():
    for i,j in z:
        print("z[" + i + ", " + j + "]=", z[i,j], sep='', end=' ')

fill()
print_field()
```

First loop over ***ti.j***, then ***ti.i***

```
import taichi as ti
ti.init(arch = ti.cpu, cpu_max_num_threads=1)

z = ti.field(ti.i32)
ti.root.dense(ti.ij, (2,2)).dense(ti.ij, (2,2)).place(z)
# a block-major hierarchical layout, size = 4x4

@ti.kernel
def fill():
    for i,j in z:
        z[i, j] = i*10 + j

@ti.kernel
def print_field():
    for i,j in z:
        print("z[" + i + ", " + j + "]=", z[i,j], sep='', end=' ')

fill()
print_field()
```

First loop over ***ti.j***, then ***ti.i***,
in 2x2 blocks

Why do we need block-major access?


ti-ink (42.99 FPS)



```
@ti.kernel
def update_flow(self):
    for P in ti.grouped(self.Flow):
        # self.Flow[P]=self.FlowNext[P]
        for i in ti.static(range(9)):
            prePos = P-self.e[i]
            self.Flow[P][i] = self.kar_avg[P][i] * \
                (self.FlowNext[P][self.k[i]] - self.FlowNext[prePos][i]) + \
                self.FlowNext[prePos][i]
```

```
self.e = ti.Vector.field(2, dtype=int, shape=9)
self.e[0] = ti.Vector([0, 0])
self.e[1] = ti.Vector([0, 1])
self.e[2] = ti.Vector([-1, 0])
self.e[3] = ti.Vector([0, -1])
self.e[4] = ti.Vector([1, 0])
self.e[5] = ti.Vector([1, 1])
self.e[6] = ti.Vector([-1, 1])
self.e[7] = ti.Vector([-1, -1])
self.e[8] = ti.Vector([1, -1])
```

Why do we need block-major access?



9-point stencil

```
self.e = ti.Vector.field(2, dtype=int, shape=9)
self.e[0] = ti.Vector([0, 0])
self.e[1] = ti.Vector([0, 1])
self.e[2] = ti.Vector([-1, 0])
self.e[3] = ti.Vector([0, -1])
self.e[4] = ti.Vector([1, 0])
self.e[5] = ti.Vector([1, 1])
self.e[6] = ti.Vector([-1, 1])
self.e[7] = ti.Vector([-1, -1])
self.e[8] = ti.Vector([1, -1])
```

Array of structures (AoS) v.s. structure of arrays (SoA) in C/C++

```
struct S1
{
    int x[8];
    int y[8];
}
S1 soa;
```

```
struct S2
{
    int x;
    int y;
}
S2 aos[8];
```



SoA



AoS

AoS v.s. SoA, which one is better?

- It really depends...

```
struct S1
{
    int x[8];
    int y[8];
}
S1 soa;

do_something(soa.x[0]);
do_something(soa.x[1]);
```

```
struct S2
{
    int x;
    int y;
}
S2 aos[8];

do_something(aos[0].x);
do_something(aos[0].y);
```



SoA



AoS

SoA in Taichi

```
x = ti.field(ti.i32)
y = ti.field(ti.i32)
ti.root.dense(ti.i, 8).place(x)
ti.root.dense(ti.i, 8).place(y)
# address: low ..... high
#           x[0]  x[1] ... x[7]  y[0]  y[1] ... y[7]
```



SoA

AoS in Taichi

```
x = ti.field(ti.i32)
y = ti.field(ti.i32)
ti.root.dense(ti.i, 8).place(x, y)
# address: low ..... high
#           x[0]  y[0]  x[1]  y[1] ... x[7]  y[7]
```



AoS

AoS in Taichi

- Only same-shaped fields can be placed in AoS fashion

```
x = ti.field(ti.i32)
y = ti.field(ti.i32)
ti.root.dense(ti.i, 8).place(x)
ti.root.dense(ti.i, 16).place(y)
# different-shaped fields x and y can not be placed in AoS fashion
```

- Shapes are determined by SNodes

```
x = ti.field(ti.i32)
y = ti.Vector.field(2, ti.i32)
ti.root.dense(ti.i, 8).dense(ti.j, 8).place(x, y)
# scalar field x and vector field y can be placed in AoS fashion
```


Switching between AoS and SoA in Taichi

```
x = ti.field(ti.i32)
y = ti.field(ti.i32)
ti.root.dense(ti.i, 8).place(x, y)
```

```
x = ti.field(ti.i32)
y = ti.field(ti.i32)
ti.root.dense(ti.i, 8).place(x)
ti.root.dense(ti.i, 8).place(y)
```

```
@ti.kernel
def foo():
    for i in x:
        do_something(x[i])

    for i in y:
        do_something(y[i])
```

```
@ti.kernel
def foo():
    for i in x:
        do_something(x[i])

    for i in y:
        do_something(y[i])
```

SoA Example, N-body:

```
pos = ti.Vector.field(2, ti.f32, N)
vel = ti.Vector.field(2, ti.f32, N)
force = ti.Vector.field(2, ti.f32, N)
```

```
...
```

```
@ti.kernel
```

```
def update():
```


```
    dt = h/substepping
```

```
    for i in range(N):
```

```
        #symplectic euler
```

```
        vel[i] += dt*force[i]/m
```

```
        pos[i] += dt*vel[i]
```




AoS Example, N-body:

```
pos = ti.Vector.field(2, ti.f32)
vel = ti.Vector.field(2, ti.f32)
force = ti.Vector.field(2, ti.f32)
ti.root.dense(ti.i, N).place(pos, vel, force)


...

@ti.kernel
def update():
    dt = h/substepping
    for i in range(N):
        #symplectic euler
        vel[i] += dt*force[i]/m
        pos[i] += dt*vel[i]
```




Loop over advanced data layouts

- Note



Loop over advanced data layouts

- Note = None
 - You can access your advanced data layouts using `struct-for(s)` as if they were your old friend *ti.field()* defined with *shape*.



Before moving to the next topic...

- Generates advanced dense data layouts using *ti.root*
 - Tree-structured: *SNode-trees*.
 - The SNode stands for “Structural Nodes”
 - All fields in Taichi are built using SNode-trees
 - *ti.root* is actual “the root” of an SNode-tree
 - $x = \text{ti.field}(\text{ti.f32}, \text{shape} = N) \iff x = \text{ti.field}(\text{ti.f32}) + \text{ti.root.dense}(\text{ti.i}, N). \text{place}(x)$
 - $\text{ti.root.dense}(\text{ti.ij}, (N, M)) \iff \text{ti.root.dense}(\text{ti.i}, N). \text{dense}(\text{ti.j}, M)$
 - You can append (multiple) dense cells to other dense cells
 - Row/col-major: $\text{ti.root.dense}(\text{ti.i}, N). \text{dense}(\text{ti.j}, M)$
 - Hierarchical layouts: $\text{ti.root.dense}(\text{ti.i}, N). \text{dense}(\text{ti.i}, M)$
 - SoA/AoS: $\text{ti.root.dense}(\text{ti.i}, N). \text{place}(x, y, z)$
 - You do not need to worry about the access of your data layouts
 - The Taichi struct-for handles it for you

Before moving to the next topic...

- Generates advanced dense data layouts using *ti.root*
 - Tree-structured: *SNode-trees*.
 - The SNode stands for “Structural Nodes”
 - All fields in Taichi are built using SNode-trees
 - *ti.root* is actual “the root” of an SNode-tree
 - $x = \text{ti.field}(\text{ti.f32}, \text{shape} = N) \iff x = \text{ti.field}(\text{ti.f32}) + \text{ti.root.dense}(\text{ti.i}, N). \text{place}(x)$
 - $\text{ti.root.dense}(\text{ti.ij}, (N, M)) \iff \text{ti.root.dense}(\text{ti.i}, N). \text{dense}(\text{ti.j}, M)$
 - You can append (multiple) dense cells to other dense cells
 - Row/col-major: $\text{ti.root.dense}(\text{ti.i}, N). \text{dense}(\text{ti.j}, M)$
 - Hierarchical layouts: $\text{ti.root.dense}(\text{ti.i}, N). \text{dense}(\text{ti.i}, M)$
 - SoA/AoS: $\text{ti.root.dense}(\text{ti.i}, N). \text{place}(x, y, z)$
 - You do not need to worry about the access of your data layouts
 - The Taichi struct-for handles it for you

Before moving to the next topic...

- Generates advanced dense data layouts using *ti.root*
 - Tree-structured: *SNode-trees*.
 - The SNode stands for “Structural Nodes”
 - All fields in Taichi are built using SNode-trees
 - *ti.root* is actual “the root” of an SNode-tree
 - $x = \text{ti.field}(\text{ti.f32}, \text{shape} = N) \iff x = \text{ti.field}(\text{ti.f32}) + \text{ti.root.dense}(\text{ti.i}, N). \text{place}(x)$
 - $\text{ti.root.dense}(\text{ti.ij}, (N, M)) \iff \text{ti.root.dense}(\text{ti.i}, N). \text{dense}(\text{ti.j}, M)$
 - You can append (multiple) dense cells to other dense cells
 - Row/col-major: $\text{ti.root.dense}(\text{ti.i}, N). \text{dense}(\text{ti.j}, M)$
 - Hierarchical layouts: $\text{ti.root.dense}(\text{ti.i}, N). \text{dense}(\text{ti.i}, M)$
 - SoA/AoS: $\text{ti.root.dense}(\text{ti.i}, N). \text{place}(x, y, z)$
 - You do not need to worry about the access of your data layouts
 - The Taichi struct-for handles it for you

Before moving to the next topic...

- Generates advanced dense data layouts using *ti.root*
 - Tree-structured: *SNode-trees*.
 - The SNode stands for “Structural Nodes”
 - All fields in Taichi are built using SNode-trees
 - *ti.root* is actual “the root” of an SNode-tree
 - $x = \text{ti.field}(\text{ti.f32}, \text{shape} = N) \iff x = \text{ti.field}(\text{ti.f32}) + \text{ti.root.dense}(\text{ti.i}, N). \text{place}(x)$
 - $\text{ti.root.dense}(\text{ti.ij}, (N, M)) \iff \text{ti.root.dense}(\text{ti.i}, N). \text{dense}(\text{ti.j}, M)$
 - You can append (multiple) dense cells to other dense cells
 - Row/col-major: $\text{ti.root.dense}(\text{ti.i}, N). \text{dense}(\text{ti.j}, M)$
 - Hierarchical layouts: $\text{ti.root.dense}(\text{ti.i}, N). \text{dense}(\text{ti.i}, M)$
 - SoA/AoS: $\text{ti.root.dense}(\text{ti.i}, N). \text{place}(x, y, z)$
 - You do not need to worry about the access of your data layouts
 - The Taichi struct-for handles it for you

Before moving to the next topic...

- Generates advanced dense data layouts using *ti.root*
 - `ti.i, ti.j, ti.k, ti.l <==> ti.axes(0), ti.axes(1), ti.axes(2), ti.axes(3)`
 - Currently Taichi supports at most 8 axes to `ti.axes(7)`
 - `ti.root.dense(ti.axes(0), 1).dense(ti.axes(1), 2).dense(ti.axes(2), 3).dense(ti.axes(3), 4).dense(ti.axes(4), 5).dense(ti.axes(5), 6).dense(ti.axes(6), 7).dense(ti.axes(7), 8).place(x)`
 - Get your Taichi updated to get the correct behavior for row/col-major fields
 - We have a new release today (10/12/2021)


Before moving to the next topic...

- Generates advanced dense data layouts using *ti.root*
 - `ti.i, ti.j, ti.k, ti.l <==> ti.axes(0), ti.axes(1), ti.axes(2), ti.axes(3)`
 - Currently Taichi supports at most 8 axes to `ti.axes(7)`
 - `ti.root.dense(ti.axes(0), 1).dense(ti.axes(1), 2).dense(ti.axes(2), 3).dense(ti.axes(3), 4).dense(ti.axes(4), 5).dense(ti.axes(5), 6).dense(ti.axes(6), 7).dense(ti.axes(7), 8).place(x)`
 - Get your Taichi updated to get the correct behavior for row/col-major fields
 - We have a new release (ver. 0.8.3) today (10/12/2021)

Sparse data layouts


The SNode-tree

- root: the root of the data structure
- dense: a fixed-length contiguous array.




The SNode-tree

- root: the root of the data structure
- dense: a fixed-length contiguous array.
- bitmasked: similar to dense, but it also uses a mask to maintain **sparsity** information, one bit per child.
- pointer: stores pointers instead of the whole structure to save memory and maintain **sparsity**



Sparse computation! but why?

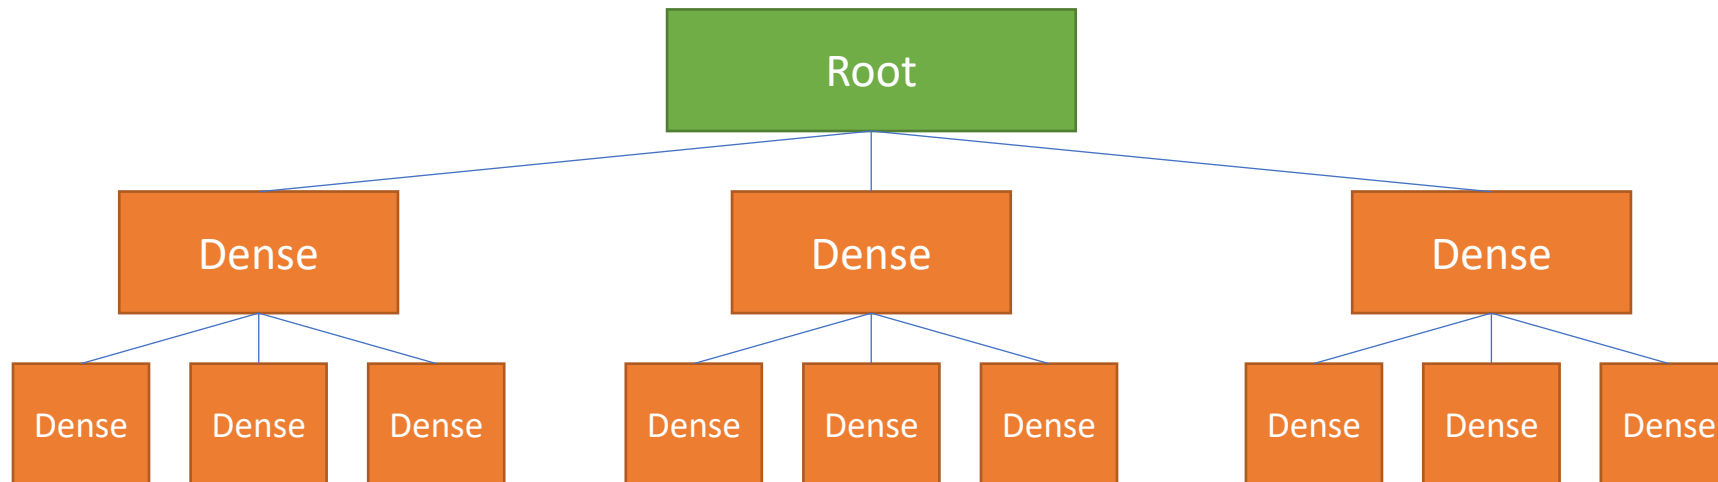
- MPM simulation →→
 - 256x256 grid cells in total
 - Subdivided to 16x16 blocks
 - Each block has 16x16 grid cells
 - Allocating memory for the total 256x256 grid cells is a waste.
 - The dark blocks are filled with zeros anyway



Sparse computation! Then how?

- A dense SNode-tree:

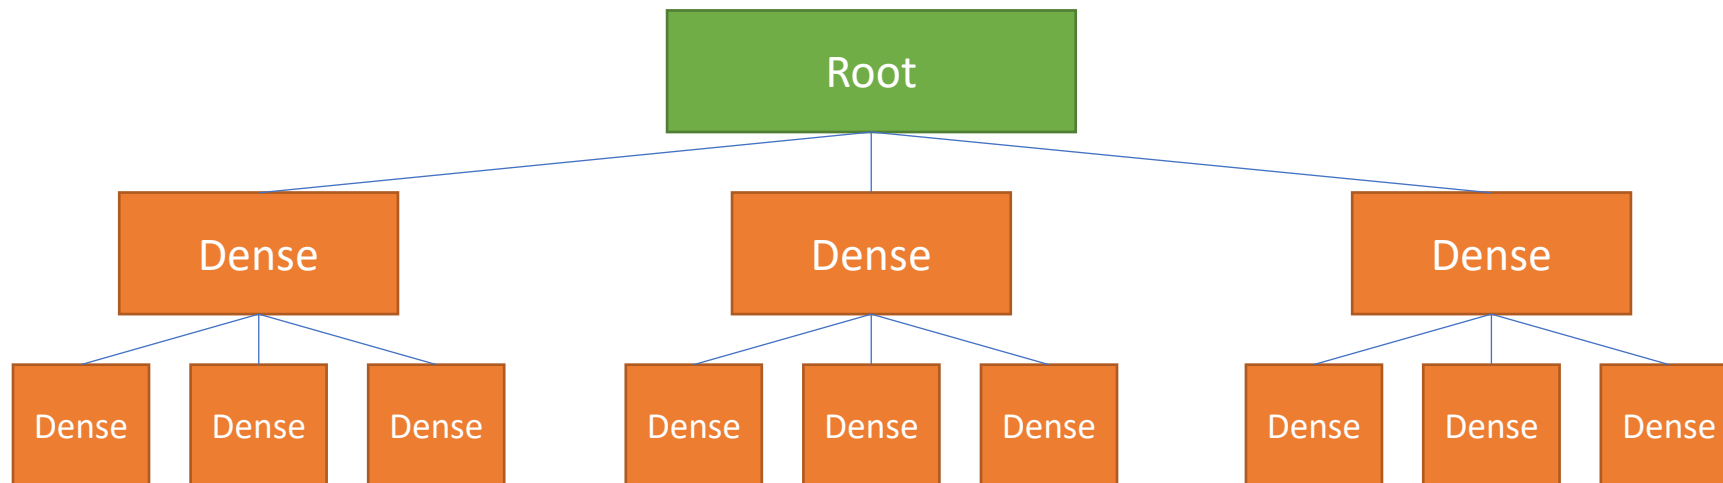
```
x = ti.field(ti.i32)
block1 = ti.root.dense(ti.i, 3)
block2 = block1.dense(ti.j, 3)
block2.place(x)
# equivalent to ti.root.dense(ti.i,3)
# .dense(ti.j,3).place(x)
```



Sparse computation! Then how?

- A dense SNode-tree:

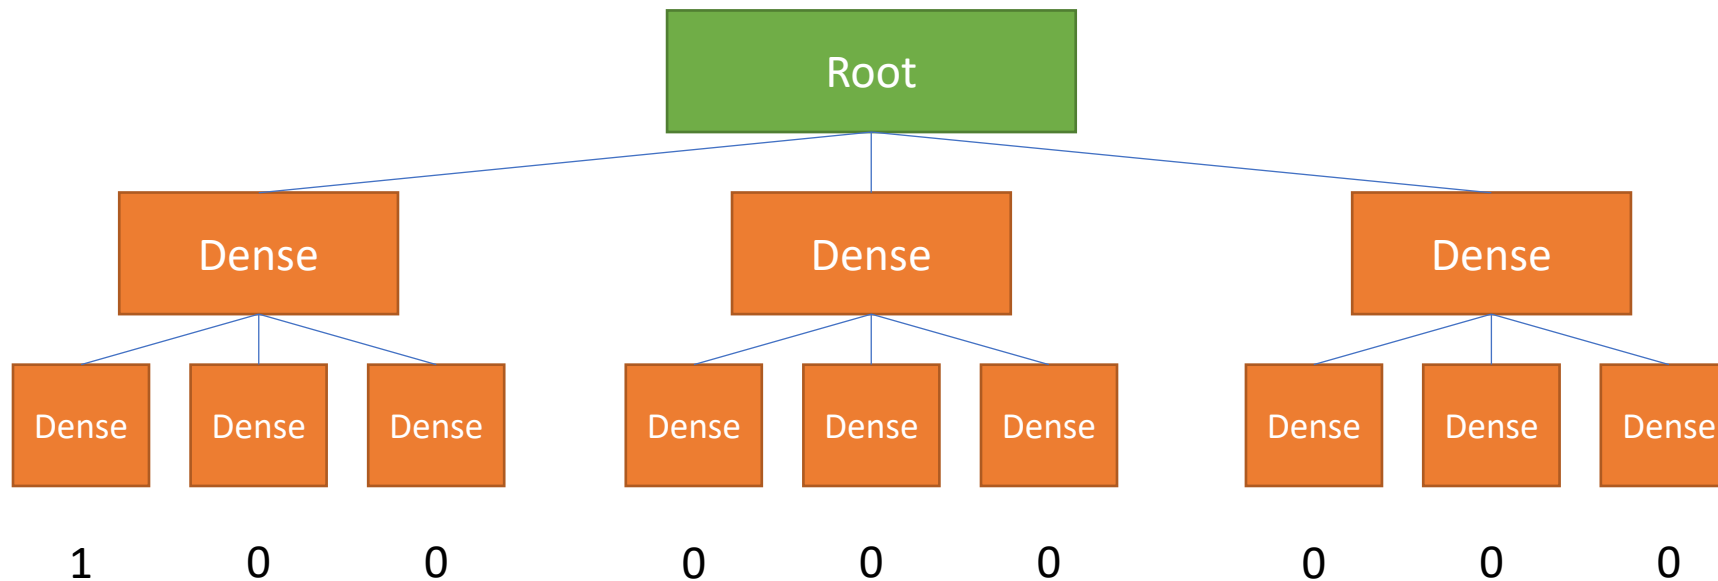
1	0	0
0	0	0
0	0	0



Sparse computation! Then how?

- A dense SNode-tree:

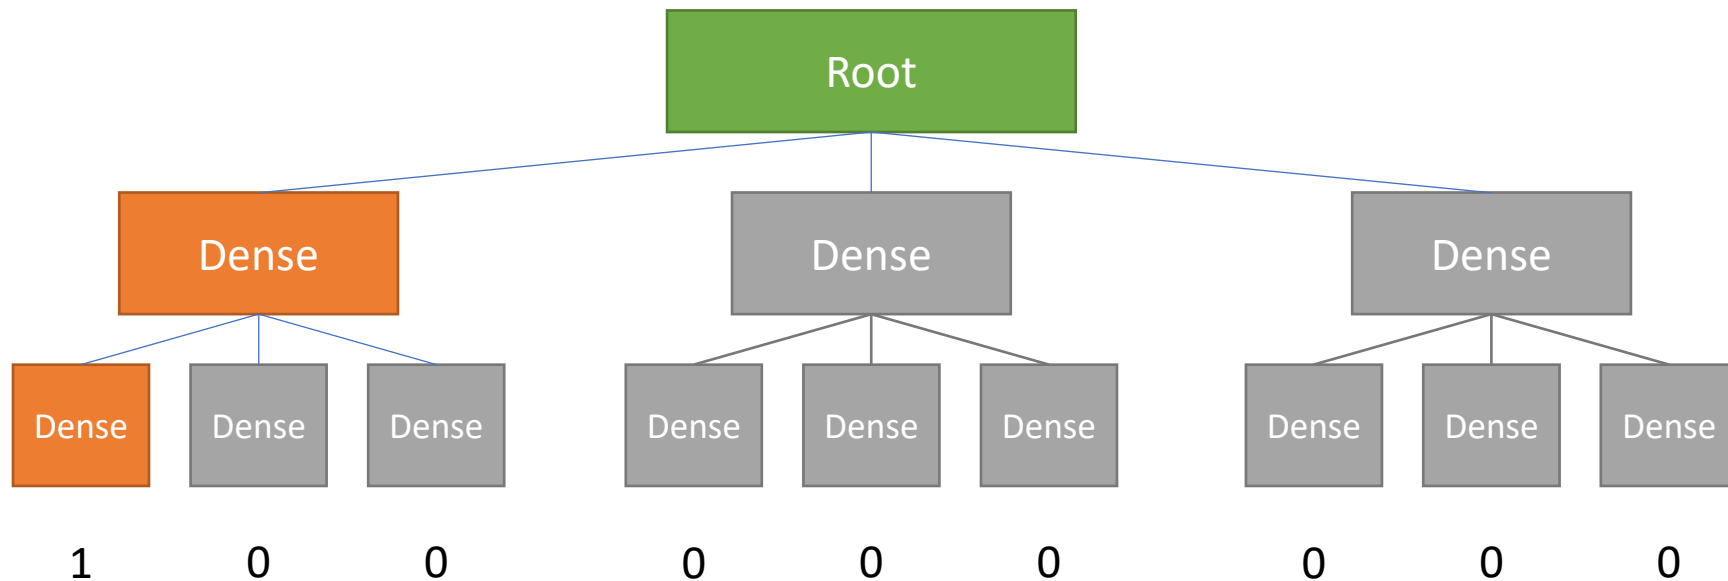
1	0	0
0	0	0
0	0	0



Sparse computation! Then how?

- A dense SNode-tree:

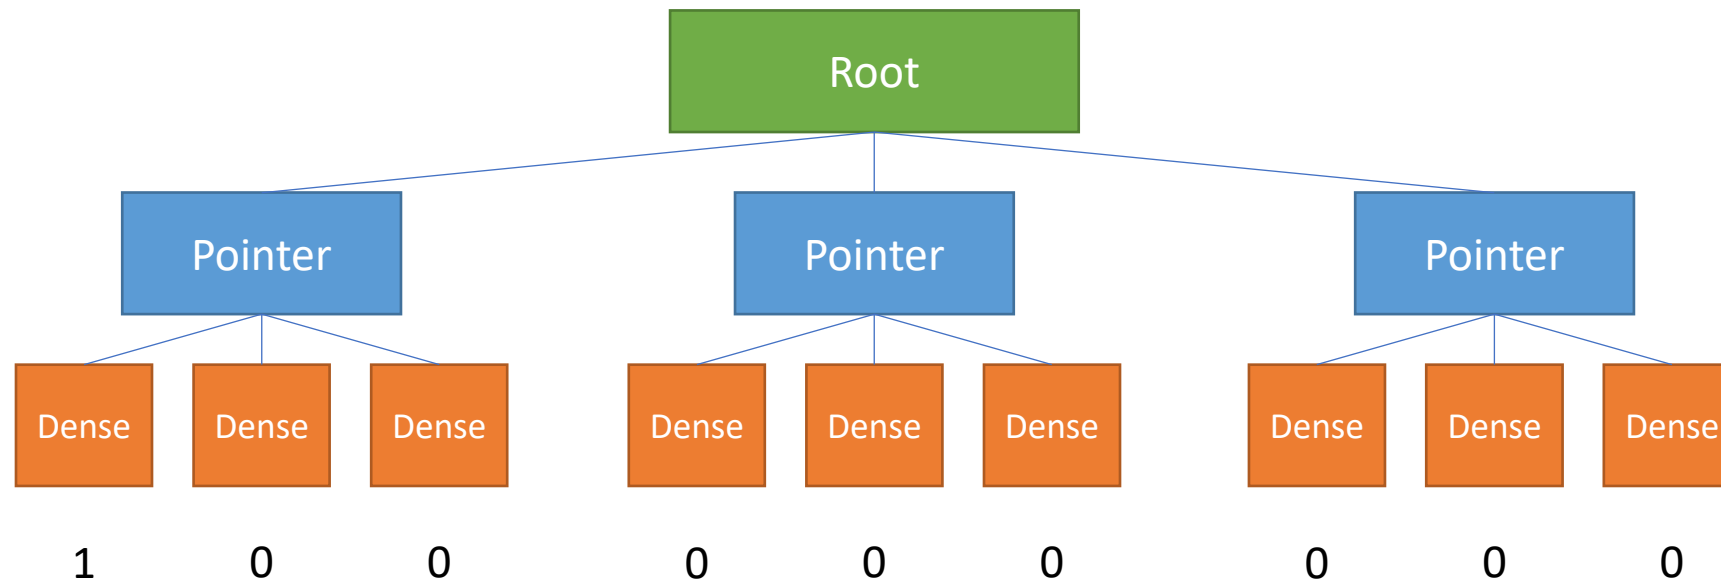
1	0	0
0	0	0
0	0	0



From .dense() to .pointer()

- A sparse SNode-tree:

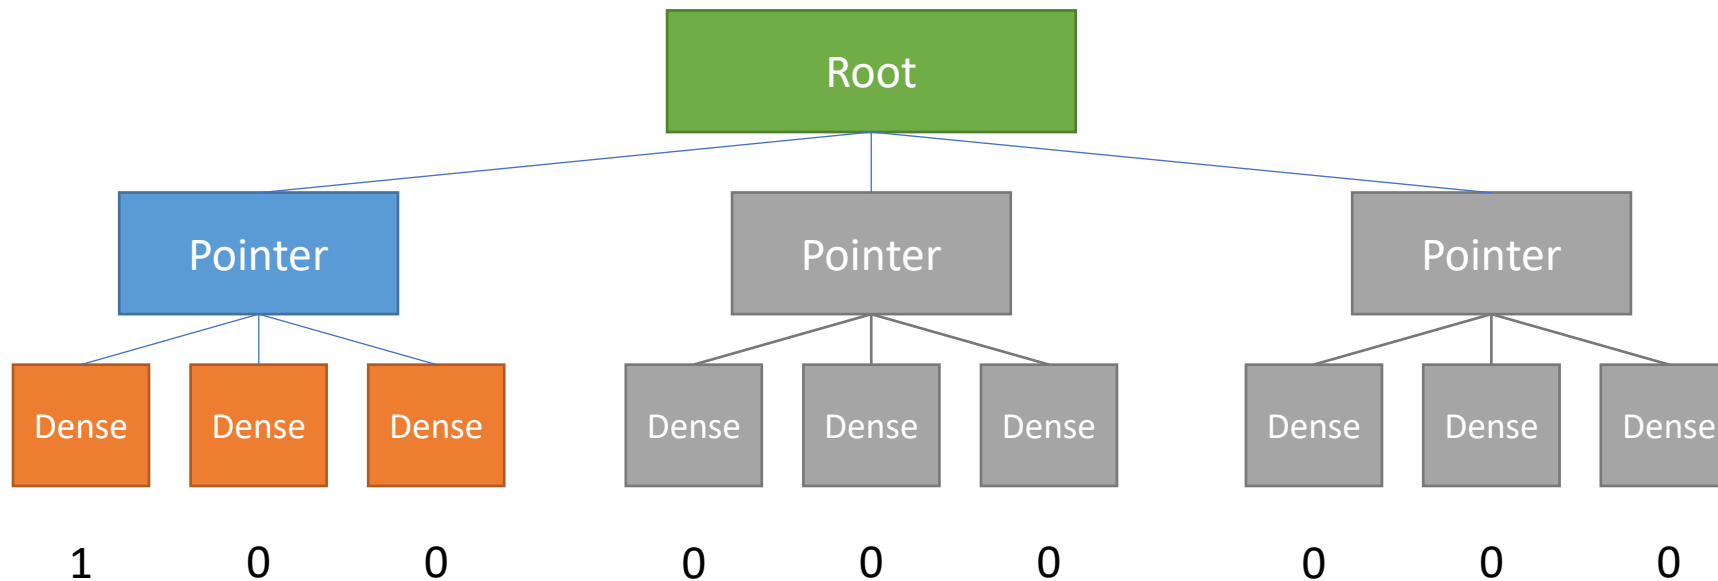
1	0	0
0	0	0
0	0	0



From .dense() to .pointer()

- A sparse SNode-tree:

1	0	0
0	0	0
0	0	0




From .dense() to .pointer()

- A sparse SNode-tree:

```
x = ti.field(ti.i32)

block1 = ti.root.pointer(ti.i, 3)
block2 = block1.dense(ti.j, 3)
block2.place(x)
# equivalent to ti.root.pointer(ti.i,3)
# .dense(ti.j,3).place(x)
```




From .dense() to .pointer()

- A sparse SNode-tree:

```
x = ti.field(ti.i32)

block1 = ti.root.pointer(ti.i, 3)
block2 = block1.dense(ti.j, 3)
block2.place(x)
# equivalent to ti.root.pointer(ti.i, 3)
# .dense(ti.j, 3).place(x)
```




Activation

- A sparse SNode-tree born empty:

```
x = ti.field(ti.i32)


block1 = ti.root.pointer(ti.i, 3)
block2 = block1.dense(ti.j, 3)
block2.place(x)
# equivalent to ti.root.pointer(ti.i,3)
# .dense(ti.j,3).place(x)
```



Activation

- Once **writing** an inactive cell:

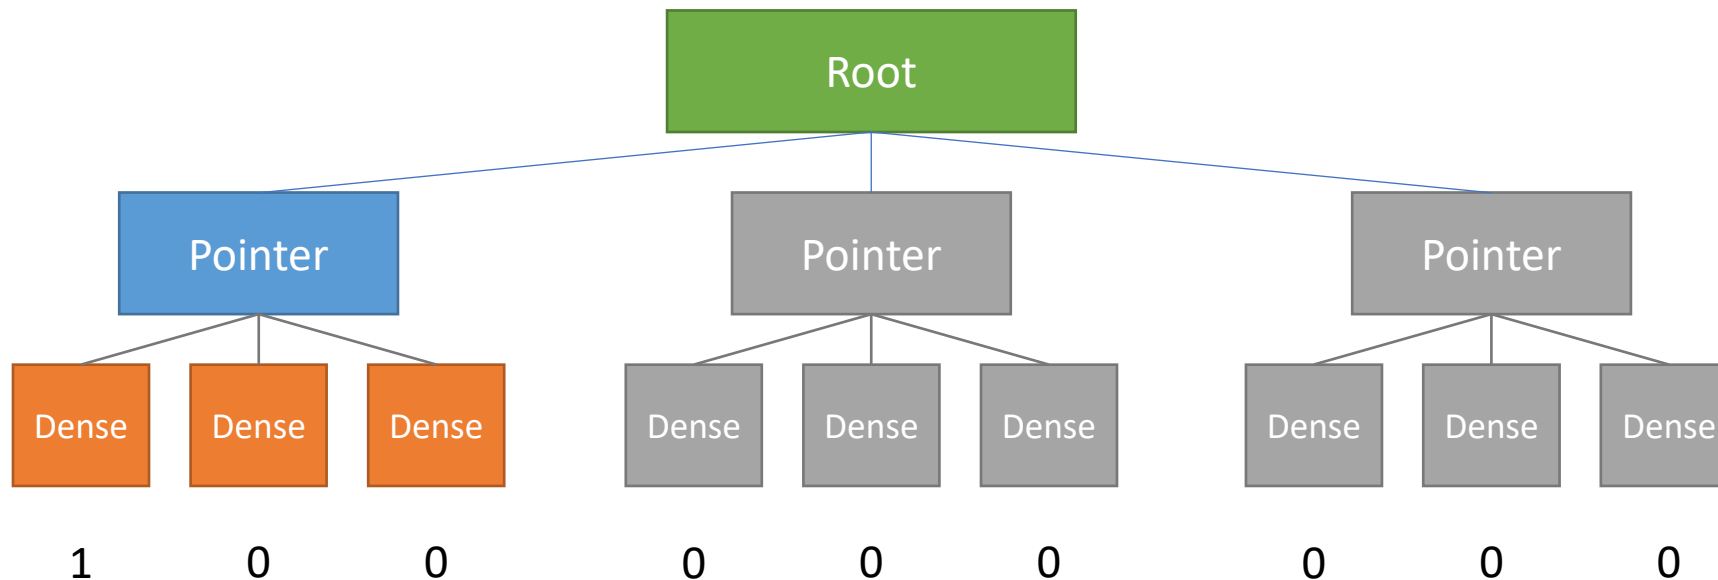
```
x[0,0] = 1  
# activates block1[0]
```



Activation

- Once **writing** an inactive cell:

```
x[0,0] = 1  
# activates block1[0] and thereby block2[0],  
block2[1] and block2[2]
```

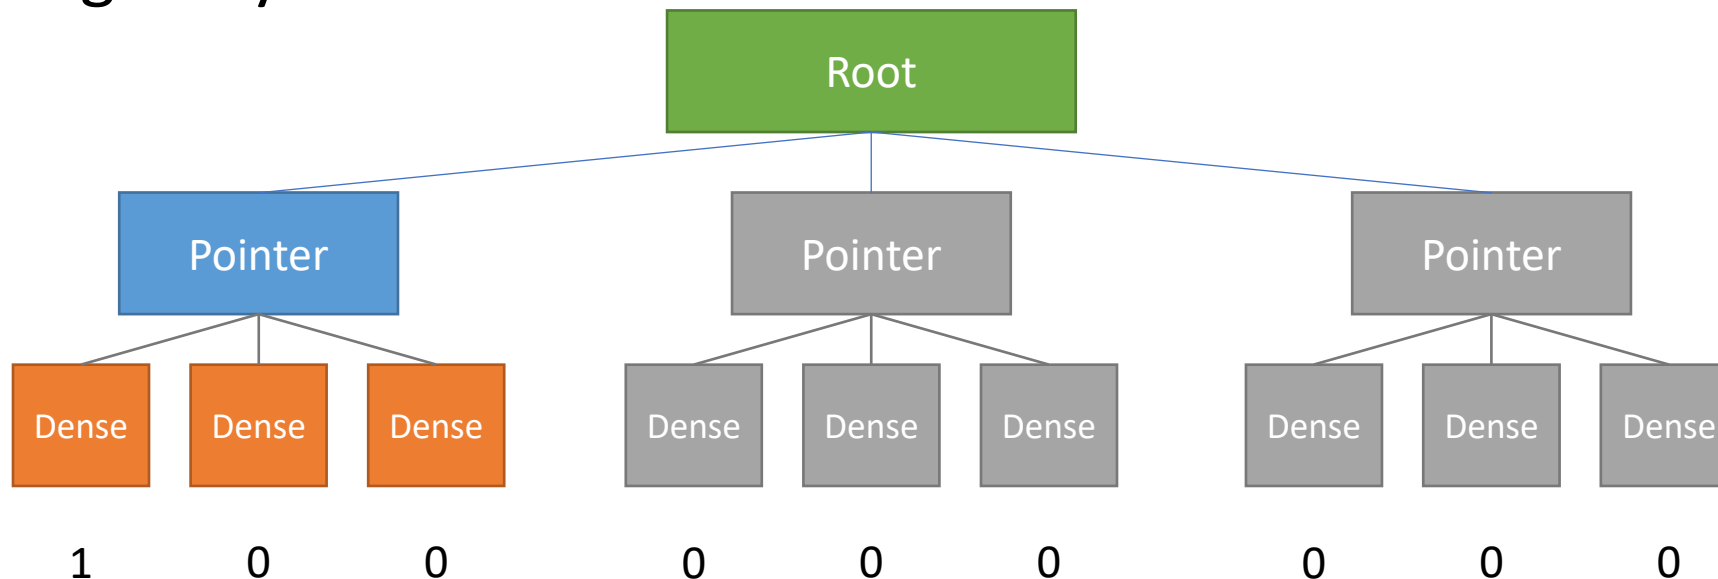


Data access in a sparse field (a sparse SNode-tree)

- Use Taichi struct-for to access a sparse field
 - Inactive pointers are skipped
- Manually accessing inactive data gives you a zero

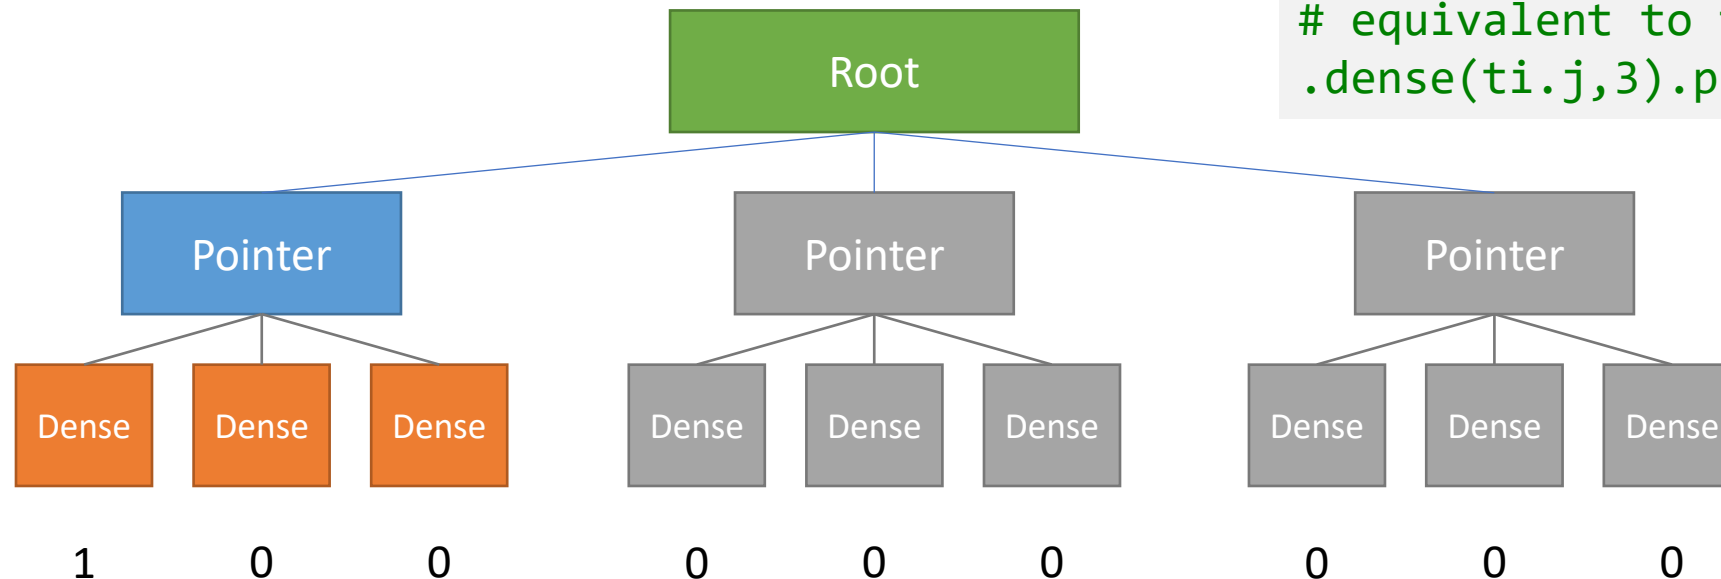
```
@ti.kernel
def access_all():
    for i,j in x:
        print(x[i, j]) # 1, 0, 0

print(x[2, 2]) # 0
```



Why activating $x[0, 1]$ and $x[0, 2]$ as well?

- Because they belong to the same dense block



```
x = ti.field(ti.i32)
```

```
block1 = ti.root.pointer(ti.i, 3)
```

```
block2 = block1.dense(ti.j, 3)
```

```
block2.place(x)
```

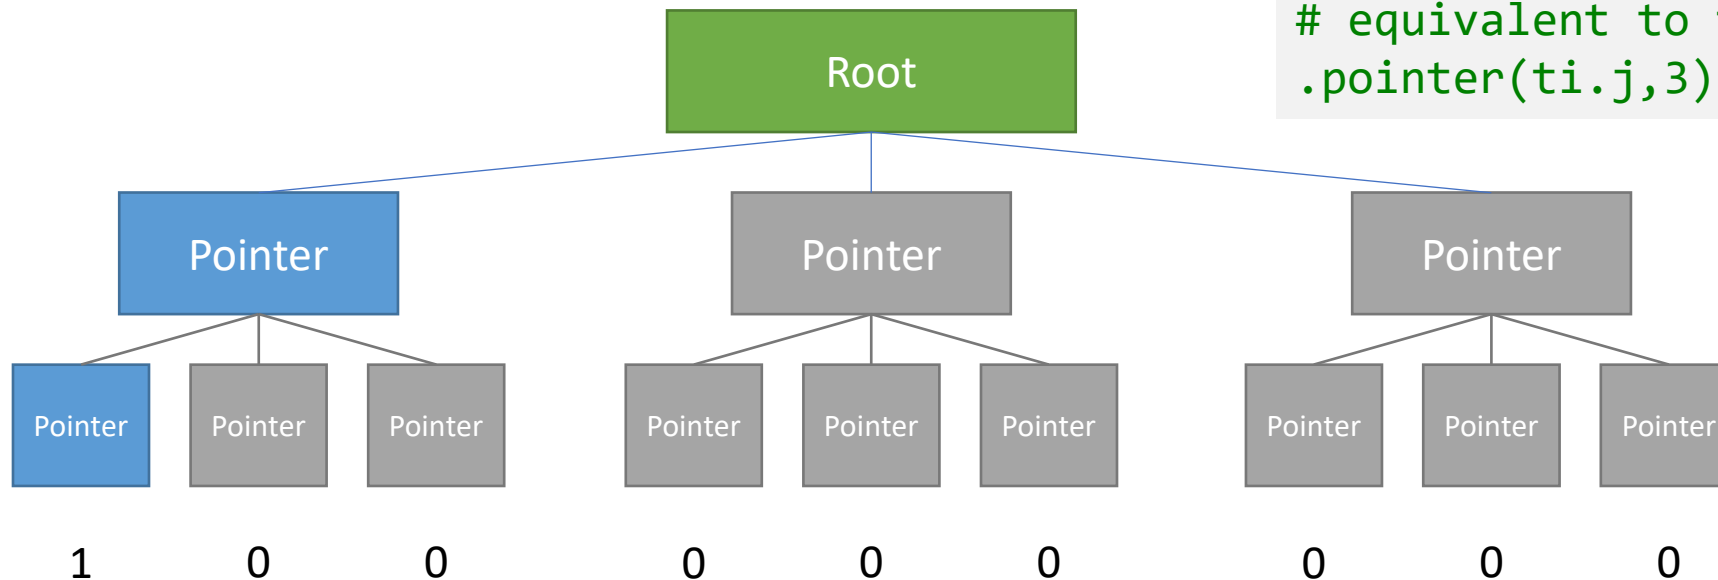
```
# equivalent to ti.root.pointer(ti.i,3)  
.dense(ti.j,3).place(x)
```

Why not using pointer everywhere?

- Bad design idea:
 - a ti.f32 \rightarrow 32 bits
 - a taichi pointer \rightarrow 64 bits

```
x = ti.field(ti.i32)

block1 = ti.root.pointer(ti.i, 3)
block2 = block1.pointer(ti.j, 3)
block2.place(x)
# equivalent to ti.root.pointer(ti.i,3)
# .pointer(ti.j,3).place(x)
```

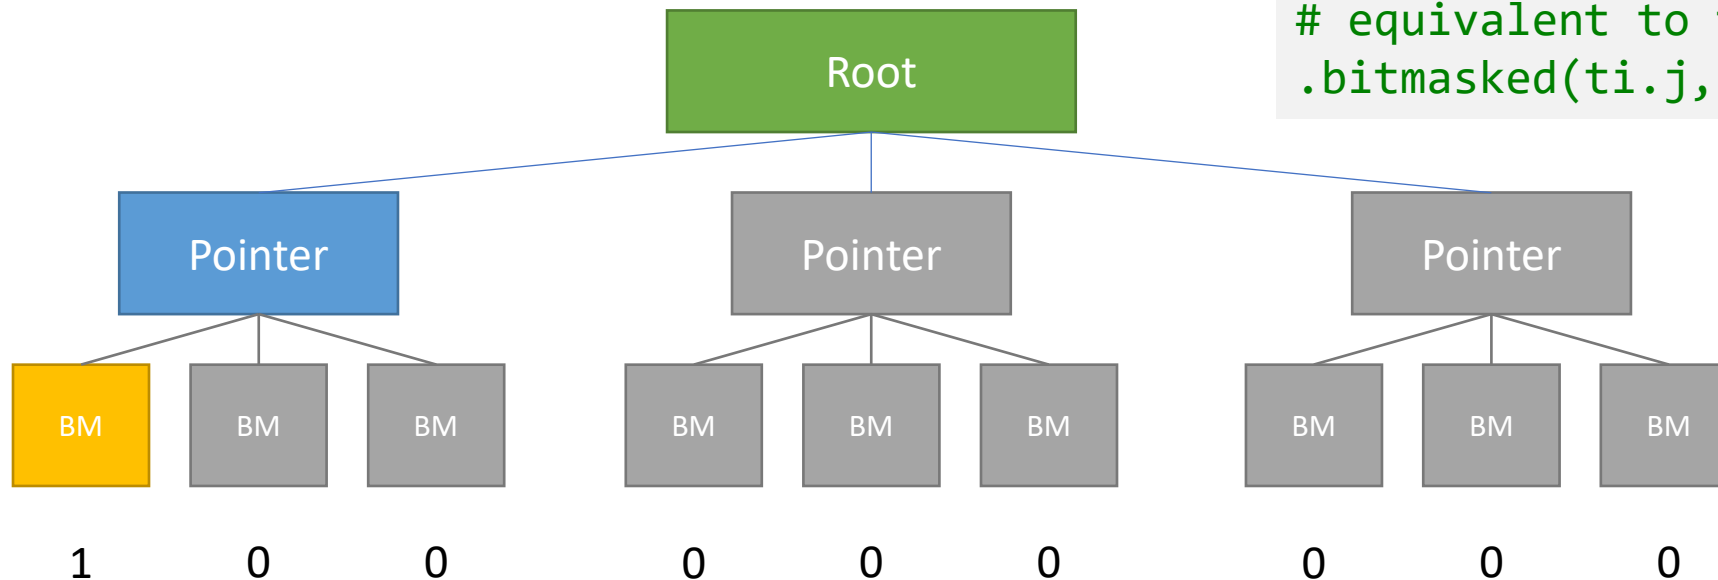


Use *bitmasks* if you really want to flag leaf cells one at a time...

- Works for leaf cells only
- Each leaf cell has its own activation flag

```
x = ti.field(ti.i32)

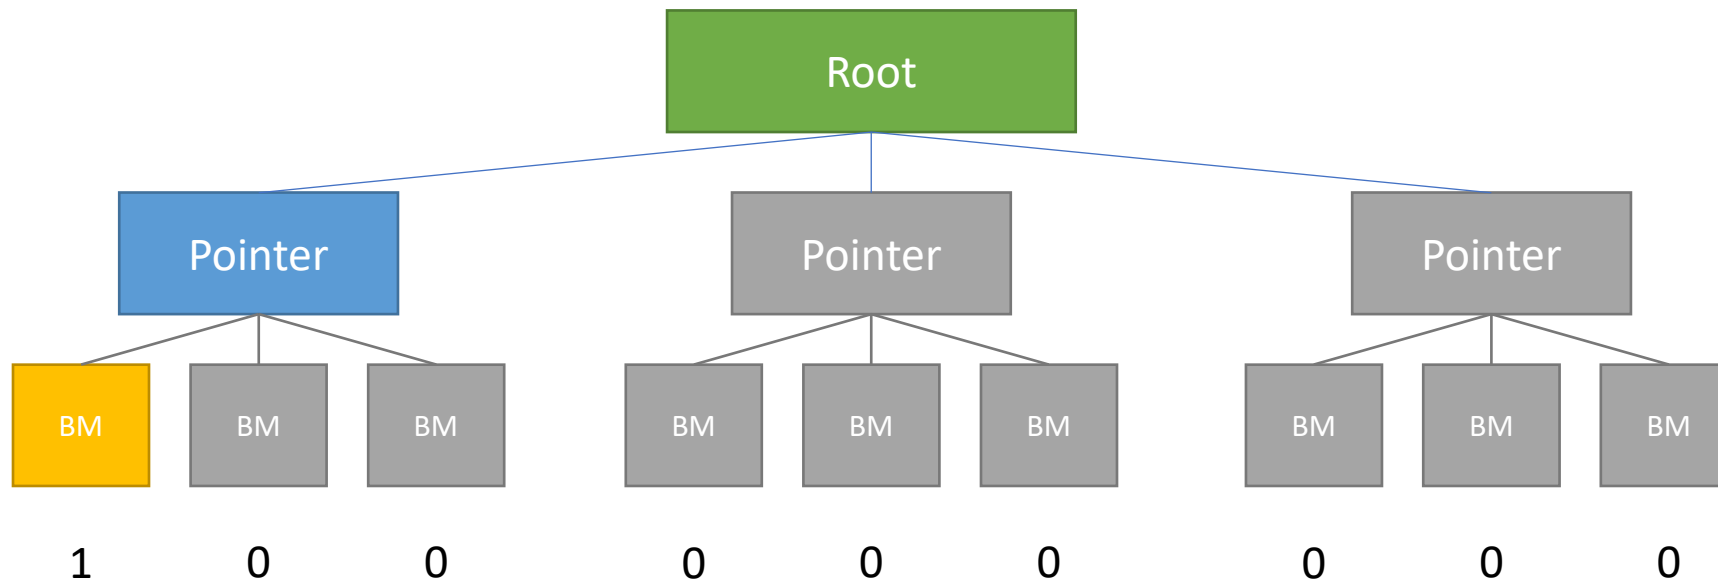
block1 = ti.root.pointer(ti.i, 3)
block2 = block1.bitmasked(ti.j, 3)
block2.place(x)
# equivalent to ti.root.pointer(ti.i,3)
# .bitmasked(ti.j,3).place(x)
```



Use *bitmasks* if you really want to flag leaf cells one at a time...


- Works for leaf cells only
- Each leaf cell has its own activation flag

```
@ti.kernel
def access_all():
    for i,j in x:
        print(x[i, j]) # 1
```



Use *bitmasks* if you really want to flag leaf cells one at a time...

- Cost 1-bit-per-cell extra
- Skip struct-for(s) when bitmasked inactive



Manual sparse field manipulation

- [API](#)

- Check activation status:

- `ti.is_active(snode, [i,j,...])`
 - for example: `ti.is_active(block1, [0])` `#=True`

- Activate/deactivate cells:

- `ti.activate/deactivate(snode, [i,j])`

- Deactivate a cell and its children:

- `snode.deactivate_all()`

- Compute the index of ancestor

- `ti.rescale_index(snode/field, ancestor_snode, index)`
 - for example: `ti.rescale_index(block2, block1, [4])` `#=1`

```
x = ti.field(ti.i32)
```

```
block1 = ti.root.pointer(ti.i, 3)
block2 = block1.dense(ti.j, 3)
block2.place(x)
# equivalent to ti.root.pointer(ti.i,3)
# .dense(ti.j,3).place(x)
```

Manual sparse field manipulation

- [API](#)

- Check activation status:
 - `ti.is_active(snode, [i,j,...])`
 - for example: `ti.is_active(block1, [0]) != True`
- Activate/deactivate cells:
 - `ti.activate/deactivate(snode, [i,j])`
- Deactivate a cell and its children:
 - `snode.deactivate_all()`
- Compute the index of ancestor
 - `ti.rescale_index(snode/field, ancestor_snode, index)`
 - for example: `ti.rescale_index(block2, block1, [4]) != 1`

```
x = ti.field(ti.i32)

block1 = ti.root.pointer(ti.i, 3)
block2 = block1.dense(ti.j, 3)
block2.place(x)
# equivalent to ti.root.pointer(ti.i,3)
# .dense(ti.j,3).place(x)
```

Manual sparse field manipulation

- [API](#)

- Check activation status:
 - `ti.is_active(snode, [i,j,...])`
 - for example: `ti.is_active(block1, [0]) != True`
- Activate/deactivate cells:
 - `ti.activate/deactivate(snode, [i,j])`
- Deactivate a cell and its children:
 - `snode.deactivate_all()`
- Compute the index of ancestor
 - `ti.rescale_index(snode/field, ancestor_snode, index)`
 - for example: `ti.rescale_index(block2, block1, [4]) != 1`

```
x = ti.field(ti.i32)

block1 = ti.root.pointer(ti.i, 3)
block2 = block1.dense(ti.j, 3)
block2.place(x)
# equivalent to ti.root.pointer(ti.i,3)
# .dense(ti.j,3).place(x)
```

Manual sparse field manipulation

- [API](#)

- Check activation status:
 - `ti.is_active(snode, [i,j,...])`
 - for example: `ti.is_active(block1, [0]) != True`
- Activate/deactivate cells:
 - `ti.activate/deactivate(snode, [i,j])`
- Deactivate a cell and its children:
 - `snode.deactivate_all()`
- Compute the index of ancestor
 - `ti.rescale_index(snode/field, ancestor_snode, index)`
 - for example: `ti.rescale_index(block2, block1, [4]) == 1`
 - Do not use `4//3` to compute the index of ancestor

```
x = ti.field(ti.i32)

block1 = ti.root.pointer(ti.i, 3)
block2 = block1.dense(ti.j, 3)
block2.place(x)
# equivalent to ti.root.pointer(ti.i,3)
# .dense(ti.j,3).place(x)
```

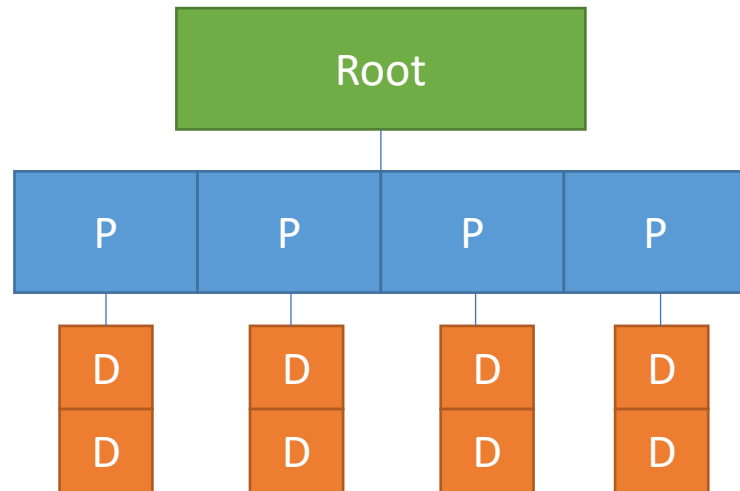
Putting things together

- Previous section:
 - Row-major v.s. col-major, flat v.s. hierarchical layouts
- This section:
 - `.dense()` v.s. `.pointer()/.bitmasked()`

Putting things together

- A column-major 2x4 2D sparse field:


```
x = ti.field(ti.i32)  
ti.root.pointer(ti.j,4).dense(ti.i,2).place(x)
```



Putting things together

- A block-majorized (block size = 3) 9x1 1D sparse field:


```
x = ti.field(ti.i32)  
ti.root.pointer(ti.i, 3).bitmasked(ti.i, 3).place(x)
```




Putting things together

- I wrote this because I could:
 - x: A column-major 2x3 2D sparse field
 - y/z: block-major sparse 4x1 1D sparse fields
 - y and z share the same sparsity pattern on p2

```
x = ti.field(ti.i32)
y = ti.field(ti.i32)
z = ti.field(ti.i32)
p1 = ti.root.pointer(ti.j, 3)
p2 = ti.root.pointer(ti.i, 2)
d11 = p1.dense(ti.i, 2)
d21 = p2.dense(ti.i, 2)
b22 = p2.bitmasked(ti.i, 2)
d11.place(x)
d21.place(y)
b22.place(z)
```



A rolling Taichi [[Code](#)]



```
n = 512
x = ti.field(ti.i32)

block1 = ti.root.pointer(ti.ij, n // 64)
block2 = block1.pointer(ti.ij, 4)
block3 = block2.pointer(ti.ij, 4)
block3.dense(ti.ij, 4).place(x)
```

The grid is divided into 8x8 *block1* containers;
Each *block1* container has 4x4 *block2* cells;
Each *block2* container has 4x4 *block3* cells;
Each *block3* container has 4x4 *pixel* cells;
Each *pixel* contains an i32 value $x[i, j]$.

Sparse data layouts

- Append more types to your SNode-tree:
 - .pointer() to represent sparse cells
 - .bitmasked() to represent sparse leaf cells
- Activate cells (and its ancestors) by writing
 - $x[0,0] = 1$
- Use Taichi struct-for(s) to access sparse fields
 - as if they were dense 😊

Sparse data layouts

- Append more types to your SNode-tree:
 - .pointer() to represent sparse cells
 - .bitmasked() to represent sparse leaf cells
- **Activate cells (and its ancestors) by writing**
 - $x[0,0] = 1$
- Use Taichi struct-for(s) to access sparse fields
 - as if they were dense 😊

Sparse data layouts

- Append more types to your SNode-tree:
 - .pointer() to represent sparse cells
 - .bitmasked() to represent sparse leaf cells
- Activate cells (and its ancestors) by writing
 - $x[0,0] = 1$
- Use Taichi struct-for(s) to access sparse fields
 - as if they were dense 😊

Sparse data layouts


- Limited backend compatibility
 - Supported by CPU/CUDA/Metal backends
- Sparse matrices are usually NOT implemented in Taichi via sparse data layouts.
 - Will cover it next week

Remark

- Advanced Data layouts for
 - Dense data structures:
 - .dense()
 - row-major v.s. col-major, hierarchical v.s. flat, AoS v.s. SoA
 - Sparse data structures:
 - .pointer() / .bitmasked()


A bigger picture

- The SNode-tree
 - root: the root of the data structure
 - dense: a fixed-length contiguous array
 - bitmasked: similar to dense, but it also uses a mask to maintain sparsity information, one bit per child
 - pointer: stores pointers instead of the whole structure to save memory and maintain sparsity



A bigger picture

- The SNode-tree
 - root: the root of the data structure
 - dense: a fixed-length contiguous array
 - bitmasked: similar to dense, but it also uses a mask to maintain sparsity information, one bit per child
 - pointer: stores pointers instead of the whole structure to save memory and maintain sparsity
 - dynamic: variable-length array, with a predefined maximum length
- Check [Yuanming's paper](#) for more details




Taichi: a **data-oriented** programming language

- Focus on data-access
 - Faster data-access \approx better performance in GPU
- Decouple the data-structures from computation
 - No need to change your code for trying different data layouts

Homework

N-body: [[Link](#)]




- Check the performance
 - SoA v.s. AoS

Perlin noise: [[Link](#)]



- Check the performance
 - Flat layout v.s. hierarchical layout

A rolling Taichi: [[Link](#)]



- Check the performance
 - Sparse (`.pointer()`) layout v.s. dense (`.dense()`) layout

Share your homework

- Could be ANYTHING you programmed using Taichi
- Help us find your homework by using [Template](#)
- Share it with your classmates at forum.taichi.graphics
 - 太极图形课作业区: <https://forum.taichi.graphics/c/homework/14>
 - Share your Taichi zoo link or your github/gitee link
 - Compile a .gif animation at your will

Gifts for the gifted

- Next check: Nov. 9th 2021

taichi-dev / taichi Public

Watch 368 Unstar 15.6k Fork 1.6k

<> Code Issues 306 Pull requests 43 Actions Security Insights

Pulse
Contributors
Community
Commits
Code frequency
Dependency graph
Network
Forks

Dependency graph

Dependencies Dependents

Repositories that depend on taichi

59 Repositories 6 Packages

litt1598 / --Galaxy	☆ 0	🍴 0
darkwuta / 2021_taichi_course_homework	☆ 4	🍴 0
FantasyVR / hw1_three_body_simulation	☆ 0	🍴 0
yuanming-hu / taichi-course-hw1	☆ 6	🍴 1
litt1598 / jubilant-enigma	☆ 0	🍴 0
taichiCourse01 / taichi_course_homework	☆ 7	🍴 2



Final tip

- Update your Taichi to 0.8.3 (released today on 10/12/2021)
 - `python -m pip install taichi --upgrade`
- Taichi is constantly evolving:
 - Raise an issue @ <https://github.com/taichi-dev/taichi> if you think you find a bug

Questions?

本次答疑：10/14

下次直播：10/19

直播回放：Bilibili 搜索「太极图形」

主页&课件：<https://github.com/taichiCourse01>